

# Building Scalable Deployments with Multiple Goroutines

GopherCon UK 2022



Sadie Freeman



**Github Repo:**  
<https://bit.ly/3C33ICS>



@Sadef  
on  
Twitter

# Agenda

- Step by step example
- Scaling for speed with Goroutines
- Scaling for load with deployments
- Case study: NFL All Day



**Coffee Anyone?**

## How do we serve:

- **as much coffee as possible**
- **to as many people as possible**
- **as fast as possible?**

Speed + Load

- **Take payment**
- **Steam milk**
- **Make espresso**

1

# No Concurrency

```
func main() {
    start := time.Now()
    PayForCoffee()
    MakeEspresso()
    SteamMilk()
    log.Printf("Coffee made, 1 customer served")

    timeTaken := time.Since(start)
    log.Printf("Took %s to serve coffee", timeTaken)
}
```

1

# No Concurrency

```
func PayForCoffee() {
    time.Sleep(2 * time.Second)
    log.Printf("Coffee paid for $")
}

func MakeEspresso() {
    time.Sleep(2 * time.Second)
    log.Printf("Espresso made ☕")
}

func SteamMilk() {
    time.Sleep(2 * time.Second)
    log.Printf("Milk steamed 🥛")
}
```

# 1 No Concurrency

go run main.go

Coffee paid for 💰

Espresso made ☕

Milk steamed 🥛

Coffee made, 1 customer served

Took 6.004392209s to serve coffee

2

# Serve more customers

```
func ServeCustomer(w http.ResponseWriter, r *http.Request) {  
    start := time.Now()  
  
    numCustomers, err := strconv.Atoi(strings.TrimPrefix(r.URL.Path, "/serve-customer/"))  
    if err != nil || numCustomers == 0 {  
        numCustomers = 1  
    }  
  
    count := 0  
    for i := 0; i < numCustomers; i++ {  
        MakeCoffee()  
        count++  
    }  
  
    timeTaken := time.Since(start)  
    log.Printf("Took %s to serve coffee to %v customer(s)", timeTaken, count)  
}
```

```
func MakeCoffee() {  
    PayForCoffee()  
    MakeEspresso()  
    SteamMilk()  
}
```

2

## Serve more customers

**go run main.go**

**curl http://localhost:8080/serve-customer/3**

```
Took 18.007806208s to serve coffee to 3 customer(s)
```

Speed

3

## Use a Goroutine

```
go MakeCoffee()
```

**Threads: Set of instructions that can be run independently**

**Concurrently: Happening at the same time**

3

## Use a Goroutine

```
Took 49.958µs to serve coffee to 3 customer(s)
```

```
Coffee paid for $
```

```
Coffee paid for $
```

```
Coffee paid for $
```

```
Espresso made ☕
```

```
Espresso made ☕
```

```
Espresso made ☕
```

```
Milk steamed 🥛
```

```
Milk steamed 🥛
```

```
Milk steamed 🥛
```

4

## Add Wait Group

```
wg := sync.WaitGroup{}  
count := 0  
for i := 0; i < numCustomers; i++ {  
    wg.Add(1)  
    go MakeCoffee(&wg)  
    count++  
}  
wg.Wait()
```

```
func MakeCoffee(wg *sync.WaitGroup) {  
    defer wg.Done()  
    PayForCoffee()  
    MakeEspresso()  
    SteamMilk()  
}
```

4

## Add Wait Group

go run main.go

curl http://localhost:8080/serve-customer/3

```
Coffee paid for 💰
Coffee paid for 💰
Coffee paid for 💰
Espresso made ☕
Espresso made ☕
Espresso made ☕
Milk steamed 🥛
Milk steamed 🥛
Milk steamed 🥛
```

Took 6.005458958s to serve coffee to 3 customer(s)

5

## Add MORE Goroutines

```
func MakeCoffee(wg *sync.WaitGroup) {  
    defer wg.Done()  
  
    newWg := sync.WaitGroup{}  
    newWg.Add(3)  
    go PayForCoffee(&newWg)  
    go MakeEspresso(&newWg)  
    go SteamMilk(&newWg)  
    newWg.Wait()  
}
```

5

## Add MORE Goroutines

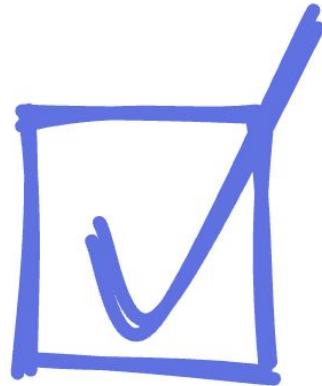
go run main.go

curl http://localhost:8080/serve-customer/3

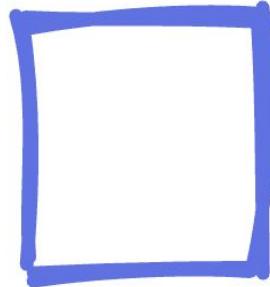
```
Coffee paid for 💵
Milk steamed 🍷
Milk steamed 🍷
Coffee paid for 💵
Espresso made ☕
Milk steamed 🍷
Coffee paid for 💵
Espresso made ☕
Espresso made ☕
Took 2.001751333s to serve coffee to 3 customer(s)
```

**MORE CUSTOMERS!!!!**





Speed



Load

## 6

# Containerize

## Dockerfile

```
FROM golang:1.18.4-alpine

WORKDIR /app

COPY go.mod .
RUN go mod download

COPY *.go .

RUN go build -o /coffee-shop

EXPOSE 8080

CMD [ "/coffee-shop" ]
```

## 6

## Containerize

## Makefile

```
build:  
    docker build --tag coffee-shop .  
  
run:    build  
    docker run -d --name coffee-shop -p 8080:8080 coffee-shop  
  
stop:  
    docker stop coffee-shop  
    docker container rm coffee-shop
```

## 6

## Containerize

```
docker build --tag coffee-shop .
[+] Building 1.2s (11/11) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 37B
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for docker.io/library/golang:1.18.4-alpine
=> [1/6] FROM docker.io/library/golang:1.18.
4-alpine@sha256:9b2a2799435823dbf6fef077a8ff7ce83ad26b382ddabf22febfc333926400e4
=> [internal] load build context
=> => transferring context: 1.33kB
=> CACHED [2/6] WORKDIR /app
=> CACHED [3/6] COPY go.mod .
=> CACHED [4/6] RUN go mod download
=> CACHED [5/6] COPY *.go .
=> CACHED [6/6] RUN go build -o /coffee-shop
=> exporting to image
=> => exporting layers
=> => writing image sha256:edb9c4e028c289b5a86f914ff8d43826313a7e37dfac0512894b4f7f251429f6
=> => naming to docker.io/library/coffee-shop
```

coffee-shop

IN USE

latest

8adda91d151f

20 minutes ago

331.3 MB



coffee-shop  
43bd5d6d7c8f

coffee-shop

Running

8080

23 minutes ago



## 6

## Containerize

<  coffee-shop coffee-shop  
RUNNING Logs

```
2022/08/09 18:37:38 Coffee paid for $  
2022/08/09 18:37:38 Milk steamed   
2022/08/09 18:37:38 Milk steamed   
2022/08/09 18:37:38 Coffee paid for $  
2022/08/09 18:37:38 Coffee paid for $  
2022/08/09 18:37:38 Milk steamed   
2022/08/09 18:37:38 Espresso made   
2022/08/09 18:37:38 Espresso made   
2022/08/09 18:37:38 Espresso made   
2022/08/09 18:37:38 Took 2.001408751s to serve coffee to 3 customer(s)
```

# Deploy on Kubernetes

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: coffee-shop
  namespace: default
spec:
  replicas: 1
  selector:
    matchLabels:
      coffee-shop: web
  template:
    metadata:
      labels:
        coffee-shop: web
    spec:
      containers:
        - name: coffee-shop
          image: coffee-shop
          imagePullPolicy: Never
```

# 7 Deploy on Kubernetes

Deployments		1 item		Namespace: default		Search Deployments...	
<input type="checkbox"/>	Name	Namespace	Pods	Replicas	Age	Conditions	<input type="checkbox"/>
<input type="checkbox"/>	coffee-shop	default	1/1	1	68s	Available Progressing	<input type="checkbox"/>

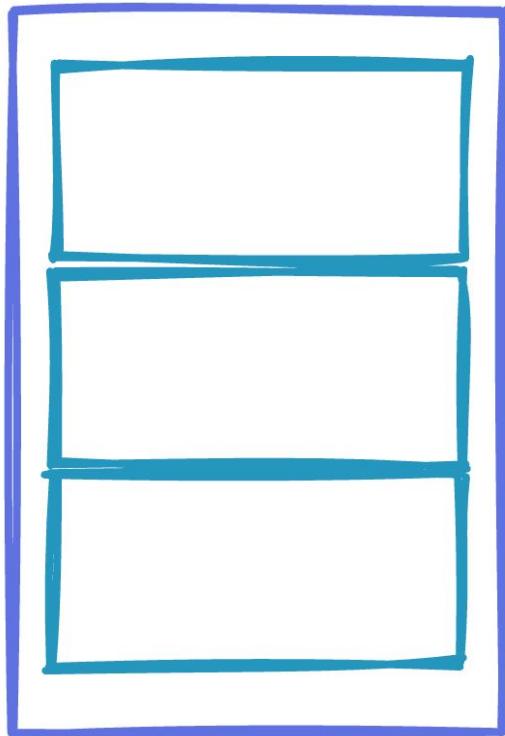
Pods									1 item		Namespace: default		Search Pods...	
<input type="checkbox"/>	Name	Namespa...	Containers	Restarts	Controlled...	Node	QoS	Age	Status	<input type="checkbox"/>	<input type="checkbox"/>			
<input type="checkbox"/>	coffee-shop-5d4756659c-c2...	default	■	0	ReplicaSet	docker-desktop	BestEffort	82s	Running	<input type="checkbox"/>	<input type="checkbox"/>			

Terminal Pod coffee-shop-5d4756659c-c2s8h X +

Namespace defa... Owner ReplicaSet coffee-shop-5d4756659c... Pod coffee-shop-5d4756659c... Container coffee-shop Search...

```
2022/08/10 20:11:10 Starting Coffee Shop Service
2022/08/10 20:12:02 Coffee paid for ☕
2022/08/10 20:12:02 Espresso made ☕
2022/08/10 20:12:02 Espresso made ☕
2022/08/10 20:12:02 Milk steamed 🥛
2022/08/10 20:12:02 Milk steamed 🥛
2022/08/10 20:12:02 Coffee paid for ☕
2022/08/10 20:12:02 Coffee paid for ☕
2022/08/10 20:12:02 Milk steamed 🥛
2022/08/10 20:12:02 Espresso made ☕
2022/08/10 20:12:02 Took 2.001860001s to serve coffee to 3 customer(s)
```

# Scaling Option 1



Vertical

## 8

# Set Resources

```
resources:  
  limits:  
    cpu: 1m  
    memory: 10Mi  
  requests:  
    cpu: 1m  
    memory: 10Mi
```

## 8 Set Resources

3

→ 2.002422959s

30

→ 2.001635459s

300

→ 3.78216496s

3000

→ OOM Killed

Restarts
1

Last Status

terminated

Reason: Reason: OOMKilled - exit code: 137

9

# Bump resources

```
resources:  
  limits:  
    cpu: 4  
    memory: 3092M  
  requests:  
    cpu: 2  
    memory: 1024M
```

9

Bump resources

30,000



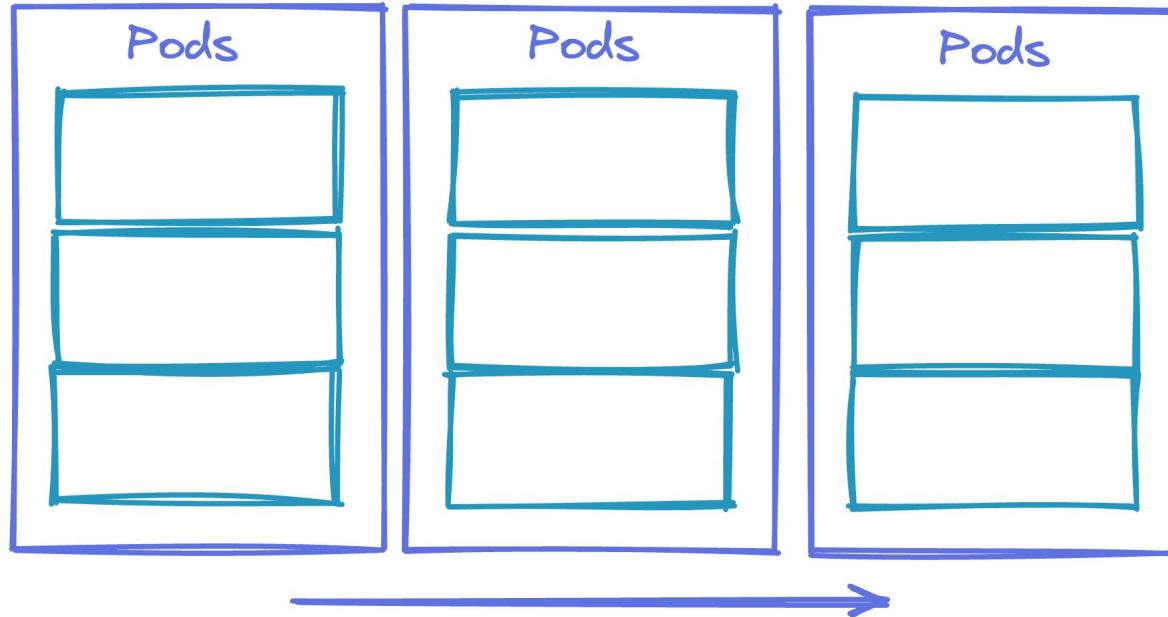
3.797142293s

3,000,000



OOM killed

# Scaling Option 2



Horizontal

10

Scale up pods

replicas: 2

10

# Scale up pods

```
apiVersion: networking.k8s.io/v1
kind: Ingress

metadata:
  namespace: default
  name: coffee-shop
  labels:
    ingress-controller: nginx
  annotations:
    kubernetes.io/ingress.class: "nginx"

spec:
  rules:
  - host: "localhost"
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: coffee-shop
            port:
              number: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: coffee-shop
  labels:
    run: coffee-shop
spec:
  ports:
  - port: 80
    protocol: TCP
    targetPort: 8080
  selector:
    run: coffee-shop
  type: NodePort
  sessionAffinity: None
```

10

# Scale up pods

Pods 2 items Namespace: default Search Pods...

Name	Namespace	Containers	Restarts	Controlled By	Node	QoS	Age	Status	⋮
coffee-shop-755cb9dc89-bsd...	default	1	0	ReplicaSet	docker-desktop	BestEffort	7m51s	Running	⋮
coffee-shop-755cb9dc89-zpp...	default	1	0	ReplicaSet	docker-desktop	BestEffort	7m51s	Running	⋮

Pod coffee-shop-755cb9dc89-bsdwp X Pod coffee-shop-755cb9dc89-zpp9w +

Namespace defa... Owner ReplicaSet coffee-shop-755cb9... Pod coffee-shop-755cb9dc89-z... Container

2022/08/16 16:10:24 Coffee paid for ☕  
2022/08/16 16:10:24 Espresso made ☕  
2022/08/16 16:10:24 Coffee paid for ☕  
2022/08/16 16:10:24 Espresso made ☕  
2022/08/16 16:10:24 Milk steamed ☕  
2022/08/16 16:10:24 Espresso made ☕  
2022/08/16 16:10:24 Coffee paid for ☕  
2022/08/16 16:10:24 Coffee paid for ☕  
2022/08/16 16:10:24 Espresso made ☕  
2022/08/16 16:10:24 Coffee paid for ☕  
2022/08/16 16:10:24 Espresso made ☕  
2022/08/16 16:10:24 Took 2.205792043s to serve coffee to 30000 customer(s)

2022/08/16 16:10:26 Espresso made ☕  
2022/08/16 16:10:26 Coffee paid for ☕  
2022/08/16 16:10:26 Espresso made ☕  
2022/08/16 16:10:26 Coffee paid for ☕  
2022/08/16 16:10:26 Espresso made ☕  
2022/08/16 16:10:26 Coffee paid for ☕  
2022/08/16 16:10:26 Espresso made ☕  
2022/08/16 16:10:26 Coffee paid for ☕  
2022/08/16 16:10:26 Espresso made ☕  
2022/08/16 16:10:26 Coffee paid for ☕  
2022/08/16 16:10:26 Espresso made ☕  
2022/08/16 16:10:26 Coffee paid for ☕  
2022/08/16 16:10:26 Espresso made ☕  
2022/08/16 16:10:26 Coffee paid for ☕  
2022/08/16 16:10:27 Took 2.265168584s to serve coffee to 30000 customer(s)

10

Scale up pods

30,000

→ 2.265168584s

30,000

→ 2.210111751s

Take Payment

2 Pods

Steam Milk

2 Pods

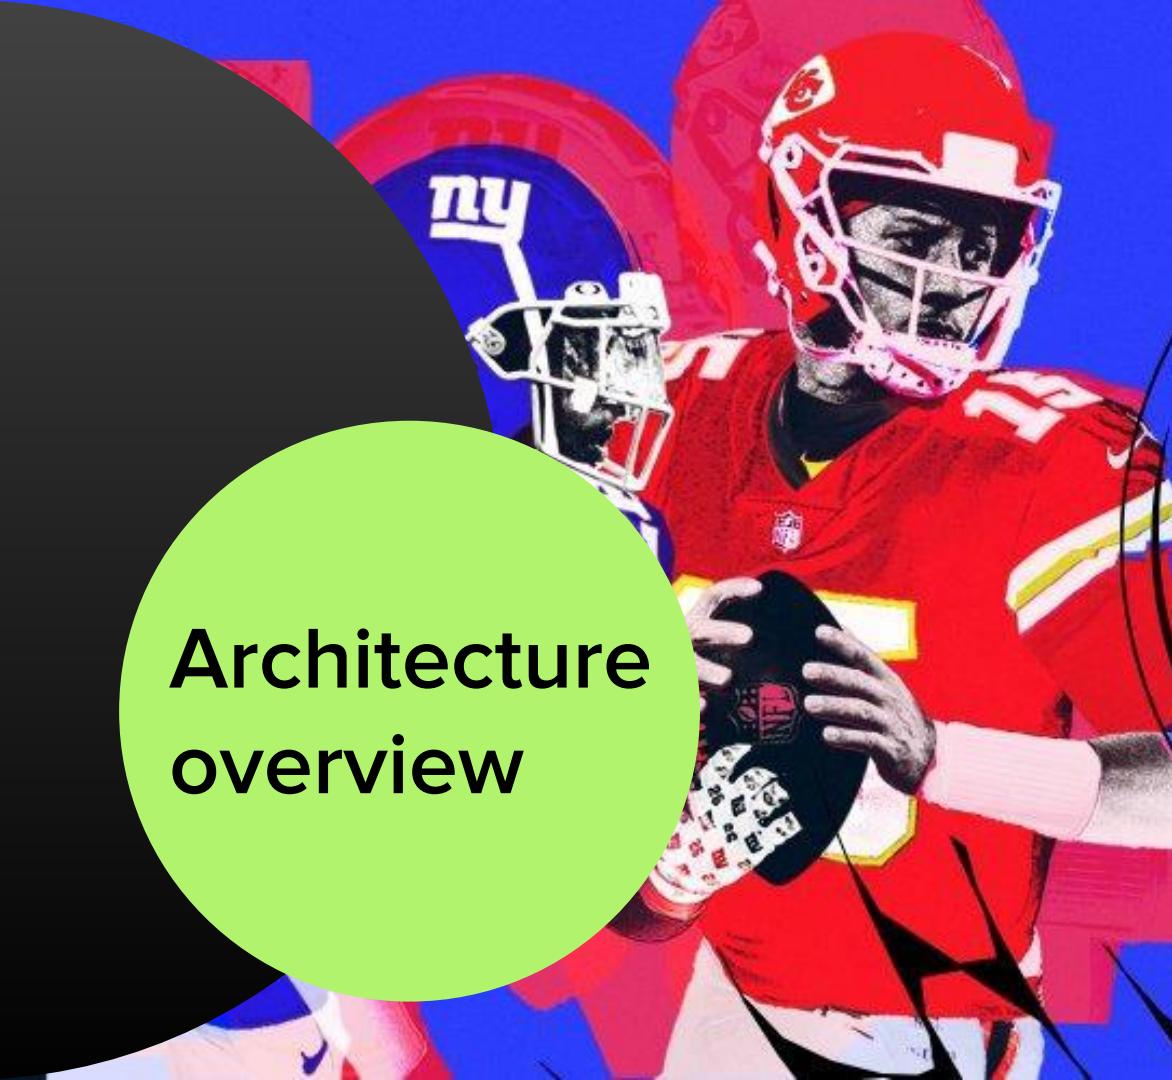
Make Espresso

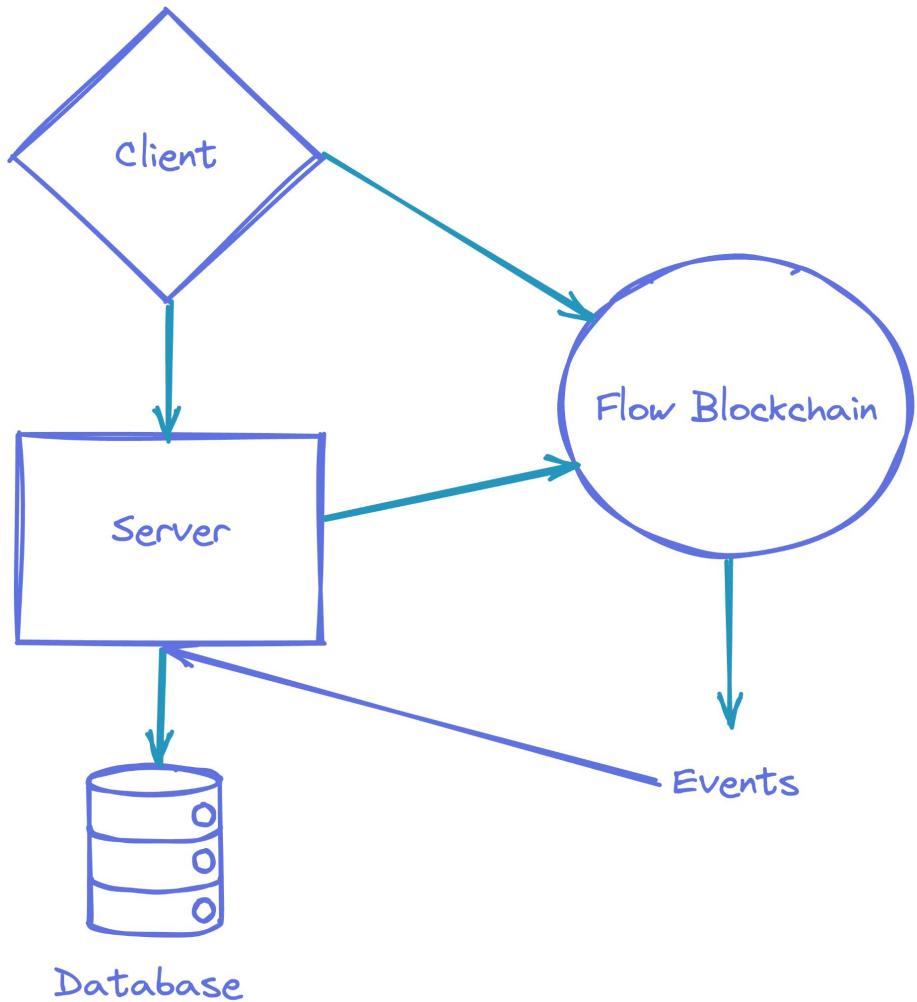
4 Pods



NFL All Day

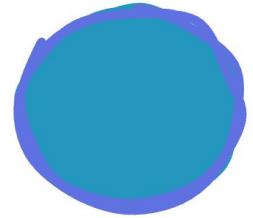
Architecture  
overview





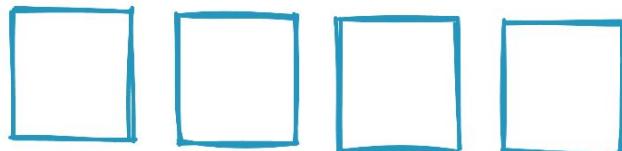
PDS -> Pack Delivery Service

- **Check Pack events**
- **Minting**
- **Sending Transactions**
- **Checking Transaction**



# No Concurrency

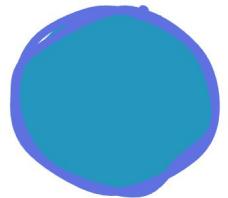
One at a time





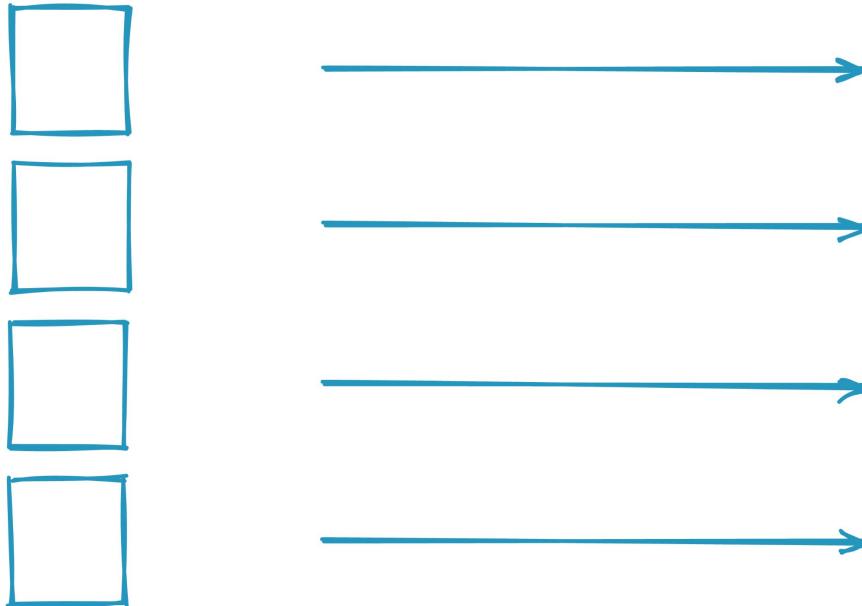
# No Concurrency

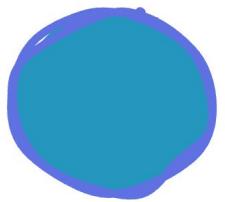
```
for {
    select {
        case <-ticker.C:
            handleMinting()
            pollCirculatingPackContractEvents()
            handleSentTransactions()
            handleSendableTransactions()
        case <-app.quit:
            cancel()
            ticker.Stop()
            return
    }
}
```



# Use Goroutines

Four at a time



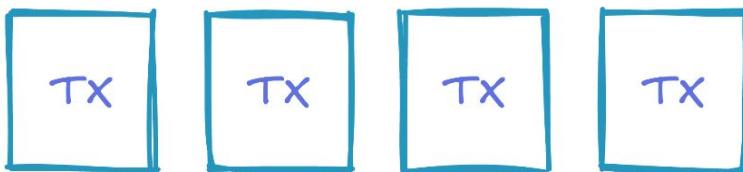


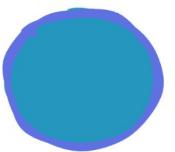
# Use Goroutines

```
go mintingPoller(app)
go packContractEventsPoller(app)
go sentTransactionsPoller(app)
go sendableTransactionPoller(app)
```

# Sendable Transactions

One at a time

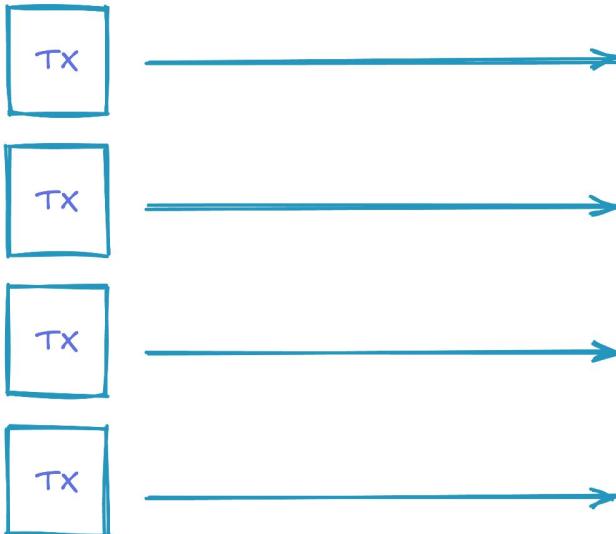


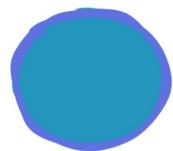


# Add More Goroutines

## Sendable Transactions

MORE at a time



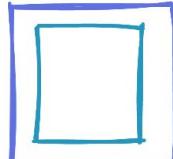


# Add More Goroutines

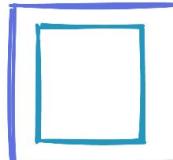
```
for _, transaction := range ts {
    wg.Add(1)
    go func(ctx context.Context, dbtx *gorm.DB, wg *sync.WaitGroup, tx *transactions.StorableTransaction) {
        defer wg.Done()
        if err := processSendableTransaction(ctx, app, logger, dbtx, tx); err != nil {
            logger.Warn("error processing storable transaction", err)
        }
    }(ctx, dbtx, &wg, &transaction)
}
wg.Wait()
```



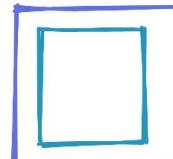
# Split Deployments



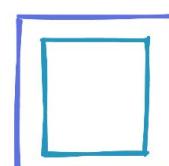
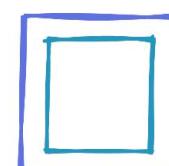
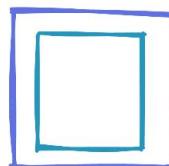
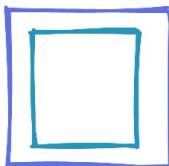
Minting



Pack Events

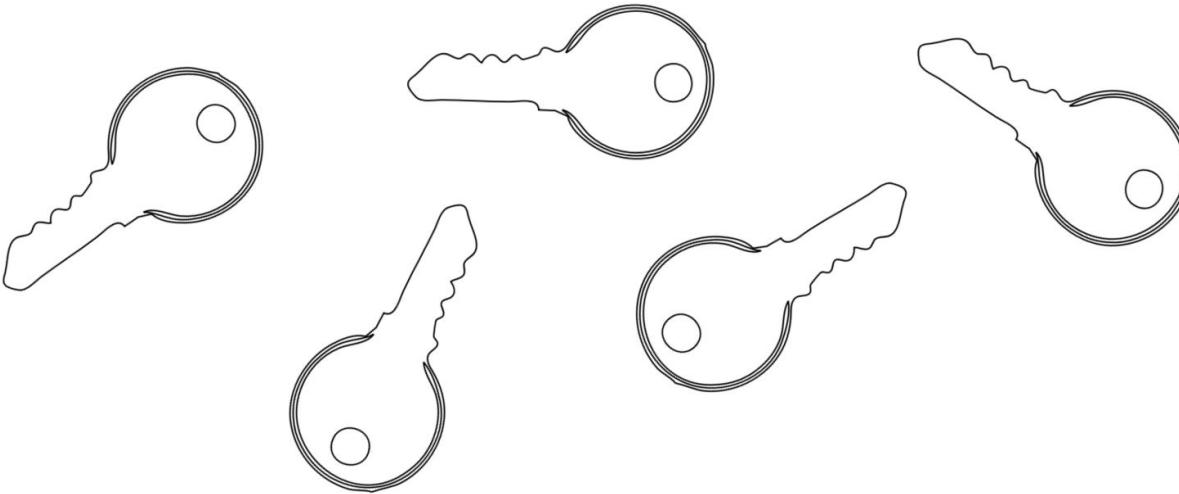


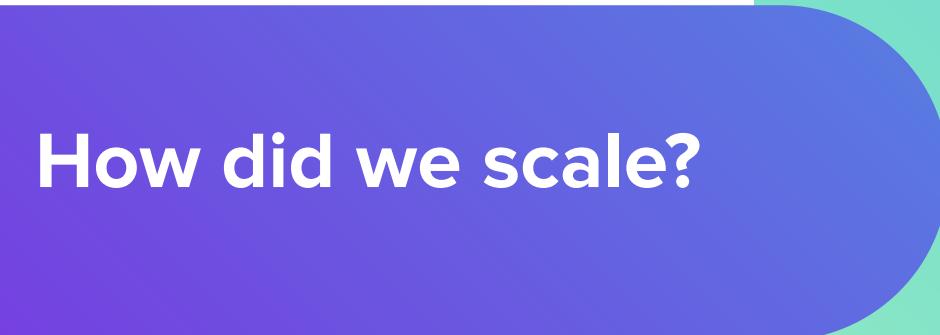
Sent Transactions



Sendable Transactions

# Keys?





**How did we scale?**

- **Use Goroutines  
(they're pretty cool)**
- **Scale horizontally +  
vertically**
- **Split different  
responsibilities in  
different  
deployments**

- Start with no concurrency
- Use Goroutines
- Use Goroutines inside of Goroutines
- Scale vertically - add resource
- Scale horizontally - add pods
- Split deployments according to Goroutine usage

# Thank You!

Eric L

Fabio K



Jude Z

Nima B

James H

Mike G

Brant H

Sean E

Ksenia O  
James H

Thilina R

Kate N

Dan B

Ryan W

Ana M

Deewai A

@Sadef  
on  
Twitter

Quinn H